

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

## Packages

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import time
import warnings
import importlib

from fastprogress import master_bar, progress_bar

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import roc_curve, roc_auc_score, classification_report
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score
from sklearn.metrics import recall_score, balanced_accuracy_score
from sklearn.metrics import average_precision_score, log_loss
from sklearn.metrics import precision_recall_curve, auc, average_precision_score
from sklearn.neural_network import MLPClassifier
from sklearn import ensemble, preprocessing
from sklearn.preprocessing import StandardScaler
from scipy.stats import ttest_rel, ttest_ind

%config InlineBackend.figure_format = 'retina'
pd.options.mode.chained_assignment = None
pd.options.display.max_columns = None
warnings.filterwarnings('ignore')
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

## Utility functions:

```

In [ ]: # instantiate a model using the parameters and import source
def get_model(model_name:str, import_module:str, model_params:dict):
    model_class = getattr(importlib.import_module(import_module), model_name)
    model = model_class(**model_params) # Instantiates the model
    return model

# returns the p_values for the each value except the best for each column
def p_stats(raw_data, raw_mean, rows, cols, per_data=False):
    # get index of algo with highest performance for each metric or dataset (in each column)
    best_algo_in_cols = raw_mean.to_numpy().argmax(axis=0)
    algo_raw = []

    for alg in raw_data:
        if per_data:
            # due to the way data is store
            algo_raw.append(np.split(alg.flatten(), cols))
        else:
            algo_raw.append(np.split(alg.flatten('F'), cols))

    p_val_test = np.ones_like(raw_mean)

    # for each metric, get the best performing algo first
    for col in range(cols):
        idx = best_algo_in_cols[col]
        # get raw data for best algo and metric(col)
        best_raw = algo_raw[idx][col]
        # run t-test between this algo(idx) and other algos
        for id_alg in range(rows):
            if (id_alg == idx):
                continue
            else:
                # run t-test between this and the best
                this_raw = algo_raw[id_alg][col]
                t_stat, p_stat = ttest_ind(best_raw, this_raw, nan_policy='omit')
                p_val_test[id_alg][col] = p_stat
    return p_val_test

def plot_roc_pr(real_y, prob_y, algo_name, data_name):
    fig, axes = plt.subplots(1,2, figsize=(12,10))
    # for roc_auc_curve
    # roc_score = roc_auc_score(real_y, prob_y)

    fpr, tpr, _ = roc_curve(real_y, prob_y)
    lab1 = 'ROC Area=%.2f' % (auc(fpr, tpr))
    axes[0].plot([1, 0], [1, 0], color='red', lw=2, linestyle='--', label='Diagonal')
    axes[0].step(fpr, tpr, label=lab1, lw=2, color='black')
    axes[0].set_title((algo_name) + ", Data: " + str(data_name))
    axes[0].set_xlabel('False Positive Rate')
    axes[0].set_ylabel('True Positive Rate')
    axes[0].legend(loc='lower left')

```

```

# set same size for two subplots
asp0 = np.diff(axes[0].get_xlim())[0] / np.diff(axes[0].get_ylim())[0]
axes[0].set_aspect(asp0)
# for precision-recall
precision, recall, _ = precision_recall_curve(real_y, prob_y)
ave_PR = average_precision_score(real_y, prob_y)
lab2 = 'PR AUC=%.2f' % (ave_PR)
# add diagonal line
axes[1].plot([0.0, 1.0], [1.0, 0.0], color='red', lw=2, linestyle='--', label='Diagonal')
axes[1].step(recall, precision, label=lab2, lw=2, color='black')
axes[1].set_title((algo_name) + ", Data: " + str(data_name))
axes[1].set_xlabel('Recall')
axes[1].set_ylabel('Precision')
axes[1].legend(loc='lower left')
asp1 = np.diff(axes[1].get_xlim())[0] / np.diff(axes[1].get_ylim())[0]
axes[1].set_aspect(asp1)
plt.show()
fig.tight_layout()

```

## Load cleaned data:

```

In [ ]: ADULT = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/cleaned/adult.csv').drop('Unnamed: 0', axis=1)
LETTER1 = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/cleaned/letter_p1.csv').drop('Unnamed: 0', axis=1)
LETTER2 = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/cleaned/letter_p2.csv').drop('Unnamed: 0', axis=1)
COVTYPE = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/cleaned/covtype5.csv').drop('Unnamed: 0', axis=1)

```

## Algorithms and their parameters:

```

In [ ]: LOGIT = {'name' : 'LogisticRegression()',
                  'name_str' : 'LogisticRegression',
                  'module' : 'sklearn.linear_model',
                  'hyperparameters' : {
                      'solver' : ['liblinear'],
                      'C' : [10**i for i in range(-8, 4)],
                  }
            }

TREES = {'name' : 'DecisionTreeClassifier()',
          'name_str' : 'DecisionTreeClassifier',
          'module' : 'sklearn.tree',
          'hyperparameters' : {
              'criterion' : ['gini', 'entropy'],
              'max_depth' : [i for i in range(1, 26, 3)],
              'min_samples_split' : [i for i in range(1, 10, 2)],
              'min_samples_leaf' : [1,3,5,7],
          }
        }

FOREST = {'name' : 'RandomForestClassifier()',
           'name_str' : 'RandomForestClassifier',
           'module' : 'sklearn.ensemble',
           'hyperparameters' : {
               'n_estimators' : [2**i for i in range(1, 11)],
               'max_features' : [1,2,4,6,8,12,16,20],
           }
         }

GBOOST = {'name' : 'GradientBoostingClassifier()',
           'name_str' : 'GradientBoostingClassifier',
           'module' : 'sklearn.ensemble',
           'hyperparameters' : {
               'n_estimators' : [2,4,8,16,32,64,128,256,512,1024,204
8],
               'max_depth' : [i for i in range(1, 10, 2)],
           }
         }

MLP_ADAM = {'name' : 'MLPClassifier()',
             'name_str' : 'MLPClassifier',
             'module' : 'sklearn.neural_network',
             'hyperparameters' : {
                 'hidden_layer_sizes' : [1,2,4,8,32,128,256],
             }
           }

NB = {'name' : 'BernoulliNB()',
      'name_str' : 'BernoulliNB',
      'module' : 'sklearn.naive_bayes',
      'hyperparameters' : {
          'alpha' : [10**i for i in range(-4, 4)]
      }}

```

# Function and variables for training classifiers

```

In [ ]: TRIALS = 5
        CV_NUM = 5

        # add all datasets/algos to a list
        data_list = [ADULT, LETTER1, LETTER2, COVTYPE]
        algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]

        # string names of variables used for printing
        dataset_names = ['ADULT', 'LETTER1', 'LETTER2', 'COVTYPE']
        algo_names = ['LOGIT', 'TREES', 'FOREST', 'GBOOST', 'MLP_ADAM', 'NB']
        metric_names = ['ACC', 'APR', 'ROC']
        scorings = {'ACC': 'accuracy', 'APR': 'average_precision', 'ROC': 'roc_auc'}

        # metrics used
        scorers = [accuracy_score, average_precision_score, roc_auc_score]

        # scores of each algorithm by data set (averaged over all metrics)
        algo_data_df = pd.DataFrame(0.0, index=algo_names, columns=dataset_names)

        # scores for each algorithm by metric (average over all data sets)
        algo_metric_df_train = pd.DataFrame(0.0, index=algo_names, columns=metric_names)
        algo_metric_df_test = pd.DataFrame(0.0, index=algo_names, columns=metric_names)

        # for each algo, store raw train/test scores for each dataset
        # number_of_dataset by trails by number_of_metrics
        raw_train = np.zeros((len(algorithms), len(dataset_names), TRIALS, len(metric_names)))
        raw_test = np.zeros((len(algorithms), len(dataset_names), TRIALS, len(metric_names)))

        # work on one algorithm
        def learn(idx_algo, algo):
            start = time.time()
            print('\n
            -----\n')
            print(f'Started training {algo["name"]}')

            # for this algo, store training metrics averaged across all datasets after the loop over data_list
            metric_across_data_train = np.zeros((len(dataset_names), len(metric_names)))

            # for this algo, store testing metrics averaged across all datasets after the loop over data_list
            metric_across_data_test = np.zeros((len(dataset_names), len(metric_names)))

            # for this algo, store raw training metrics
            # for each dataset, there is 3 metrics. For each metric, there is 5 trials

```

```

raw_metric_data_train = np.zeros((len(dataset_names), TRIALS, len(
(metric_names)))

    # for this algo, store raw training metrics
    raw_metric_data_test = np.zeros((len(dataset_names), TRIALS, len(m
etric_names)))

    # loop over all datasets
    for idx_data, data in enumerate(progress_bar(data_list)):
        start_data = time.time()
        print(f'Started on {dataset_names[idx_data]} dataset')

        # for each algo/data combo, store testing metrics averaged acr
oss 5 trials
        metric_across_trials_test = np.zeros((TRIALS, len(metric_name
s)))
        # for each algo/data combo, store training metrics averaged ac
ross 5 trials
        metric_across_trials_train = np.zeros((TRIALS, len(metric_name
s)))

        # data for precision-recall curve
        auc_real_y = []
        auc_prob_y = []

        # loop over all trials
        for trial in progress_bar(range(TRIALS)):
            start_trial = time.time()
            print("Started trial: ", trial+1)

            # pick 5000 samples for training
            X = data.drop('y', axis=1)
            Y = data['y']
            X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
train_size=5000, random_state=trial)

            # only scale data for MLP's or LOGREG
            if (algo['name_str'] in ['MLPClassifier', 'LogisticRegress
ion']):
                scaler = StandardScaler()
                X_train = pd.DataFrame(scaler.fit_transform(X_train),
columns = X.columns)
                X_test = pd.DataFrame(scaler.transform(X_test), column
s = X.columns)

                clf = eval(algo['name'])
                param = algo['hyperparameters']

                # Get each parameter that has best performance on validati
on set

                CV = GridSearchCV(clf, param, cv=CV_NUM, n_jobs=-1, scorin
g=scorings, refit='ACC')
                CV.fit(X_train, Y_train)
                result_cv = CV.cv_results_

```

```

print("\nBest param for ACC:", CV.best_params_)

# get parameters for best models for each metric
param_list = []
for metric in metric_names:
    best_id = pd.Series(result_cv['rank_test_'+str(metric)]).idxmin()
    param_list.append(result_cv['params'][best_id])

# Train n models using the 5000 samples and each of the n
best parameters
# and test on test set
clf_name = algo['name_str']
module = algo['module']
train_metrics = []
test_metrics = []

for i in range(len(param_list)):
    clf_best = get_model(clf_name, module, param_list[i])
    clf_best.fit(X_train, Y_train)

    # for roc_auc and PR curves
    pred_proba = clf_best.predict_proba(X_test)
    auc_real_y.append(Y_test)
    auc_prob_y.append(pred_proba[:,1])

    X_train_pred = clf_best.predict(X_train)
    X_test_pred = clf_best.predict(X_test)
    X_train_pred_prob = clf_best.predict_proba(X_train)
    X_test_pred_prob = clf_best.predict_proba(X_test)

    # get appropriate scoring function
    scorer = scorers[i]

    # if metric is average precision, need y-score or prob
a
    if (i==1):
        train_metrics.append(scorer(Y_train, X_train_pred_
prob[:, 1]))
        test_metrics.append(scorer(Y_test, X_test_pred_pro
b[:, 1]))
    # if scorer is roc_auc, need proba
    elif (i==2):
        train_metrics.append(scorer(Y_train, X_train_pred_
prob[:, 1]))
        test_metrics.append(scorer(Y_test, X_test_pred_pro
b[:, 1]))
    else:
        train_metrics.append(scorer(Y_train, X_train_pre
d))
        test_metrics.append(scorer(Y_test, X_test_pred))

# update the row in metric_across_trials_train
metric_across_trials_train[trial] = train_metrics

```



```

        # update the row in metric_across_trials
        metric_across_trials_test[trial] = test_metrics

        finish_trial = time.time()
        print(f'Ended trial {trial+1} in {(finish_trial - start_trial):.3f} seconds')
        print('\n')

        # plot ROC and PR curves
        auc_real_y = np.concatenate(auc_real_y)
        auc_prob_y = np.concatenate(auc_prob_y)
        plot_roc_pr(auc_real_y, auc_prob_y, algo['name_str'], dataset_names[idx_data])

        # add 5 trails by 3 metrics data to raw list
        raw_metric_data_train[idx_data] = metric_across_trials_train
        raw_metric_data_test[idx_data] = metric_across_trials_test

        # mean of metrics across trials
        mean_across_trials_train = np.mean(metric_across_trials_train, axis=0)
        mean_across_trials_test = np.mean(metric_across_trials_test, axis=0)

        # update algo-data combo with mean of mean_across_trials
        mean_algo_data = np.mean(mean_across_trials_test)
        algo_data_df.iat[idx_algo, idx_data] = mean_algo_data

        # update metric_across_data
        metric_across_data_train[idx_data] = mean_across_trials_train
        metric_across_data_test[idx_data] = mean_across_trials_test

        finish_data = time.time()
        print(f'Ended {dataset_names[idx_data]} in {(finish_data - start_data):.3f} seconds')

        # update raw_train and raw_test
        raw_train[idx_algo] = raw_metric_data_train
        raw_test[idx_algo] = raw_metric_data_test

        # mean of metrics across data
        mean_across_data_train = np.mean(metric_across_data_train, axis=0)
        mean_across_data_test = np.mean(metric_across_data_test, axis=0)
        algo_metric_df_train.iloc[idx_algo] = mean_across_data_train
        algo_metric_df_test.iloc[idx_algo] = mean_across_data_test

        finish = time.time()
        print(f'Ended {algo["name"]} in {(finish - start):.3f} seconds')
        print('\n')
    -----\n')

```

## Training Per Algorithm

## Logit

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(0, algorithms[0])
```

## Decision Trees

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(1, algorithms[1])
```

## Random Forest

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(2, algorithms[2])
```

## Gradient Boosting Classifier

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(3, algorithms[3])
```

## MLP with ADAM

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(4, algorithms[4])
```

## Naive Bayes

```
In [ ]: # algorithms = [LOGIT, TREES, FOREST, GBOOST, MLP_ADAM, NB]
learn(5, algorithms[5])
```

## Tables and Output

```
In [ ]: p_val_test_metric = p_stats(raw_test, algo_metric_df_test, len(algo_names), len(metric_names), per_data=False)
print('p_val_test_metric: for Table 2')
print(p_val_test_metric);print()

p_val_train_metric = p_stats(raw_train, algo_metric_df_train, len(algo_names), len(metric_names), per_data=False)
print('p_val_train_metric:')
print(p_val_train_metric);print()

p_val_test_data = p_stats(raw_test, algo_data_df, len(algo_names), len(data_list), per_data=True)
print('p_val_test_data: for Table 3')
print(p_val_test_data);print()
```